

Prior Work Communication: Proof-of-Control v1.0 Public Comment Filings

Custody record, v1. September 23, 2026. Filed by Basil C. Puglisi, MPA, basilpuglisi.com, under GitHub account basilpuglisi.

Standard: Proof-of-Control v1.0 working draft, Advanced AI Society, hosted as a lab of LF Decentralized Trust. Repository: github.com/LFDT-ProofOfControl/ov-poc-standard (head commit 22c7b62 at filing). Public comment window: August 1 to October 30, 2026. First contributors meeting: October 7, 2026.

What this record shows. Nine public comments were filed as GitHub issues #67 through #75 on September 23, 2026. Each comment names a gap in the draft and cites dated published work as reference under its existing license, not as contributed text. The screenshots that follow were captured by the filer at the time of filing and show each issue open under the basilpuglisi account. A SHA-256 digest of each source screenshot is listed in the exhibit index.

Filing log

Issue	Comment	Title as filed	Dated work cited (earliest first)
#67	1	C8 open item 2: clarifying human authority at Tier 4 and the root of trust	CBG (Sept 23, 2025); HAIA-RECCLIN governance architecture specification v1.7 (Feb 3, 2026); Why Agentic AI Was Always Going to Fail (June 4, 2026)
#68	2	Clarifying the relationship between open verification and human responsibility	CBG (Sept 23, 2025); HAIA-RECCLIN specification v1.7 (Feb 3, 2026); Verified AI Inference Standards Act (Mar 6, 2026)
#69	4	Clarifying principal identity and accountable ownership	CBG (Sept 23, 2025); Why Agentic AI Was Always Going to Fail (June 4, 2026)
#70	6	Source custody for research-driven requirements and crosswalk citations	SCOPE (June 2, 2026); Fault-Based Publication Ethics (June 3, 2026), cited as complementary work
#71	3	Approval fatigue: possible pattern-level signals for partial coverage	CBG (Sept 23, 2025); HAIA-RECCLIN agent architecture specification v2.2, Appendix B (Feb 4, 2026)
#72	7	Issue 13: anchoring intervals, history-rewrite risk, and halt authority	GOPEL post-quantum cryptographic agility amendment v1.2 (Mar 6, 2026); Overwatch (May 10, 2026)
#73	5	Evidence-basis declarations and classification of human decisions	CBG (Sept 23, 2025); HAIA-CARCS v1.33 (Apr 23, 2026)
#74	8	H.R. 10362 alignment: per-transaction inference evidence and a NIST-facing crosswalk	Verified AI Inference Standards Act (Mar 6, 2026)
#75	9	Recording who accepted the control set that evidence binds to	CBG (Sept 23, 2025); HAIA-RECCLIN specification v1.7 (Feb 3, 2026)

Comment 5 was filed after Comment 7, so its issue number (#73) follows #72. Comment numbers follow the filing set; issue numbers follow GitHub order.

Exhibit index

Exhibit	Content	Source file	SHA-256
A	Prefilled issue form for Comment 1 before submission, showing the title and body as drafted.	1790191857147_image.png	92aeec0d541b2f9c35f3cff141e64b84df6ce4f00e6a46e13808381a10070ed
67	Issue #67 (Comment 1), "C8 open item 2: clarifying human authority at Tier 4 and the root of trust," open under basilpuglisi.	1790191921653_image.png	8f2b864cabbc1a926a56fc45a7ce57f89f0d685fc74faf2ebad7c72cd57a6b5f
68	Issue #68 (Comment 2), "Clarifying the relationship between open verification and human responsibility," open under basilpuglisi.	1790192024034_image.png	0f25a663228ccde05d1bd92baee602c6fbbc2d26e4379357bd5ba9eb3326b674
69	Issue #69 (Comment 4), "Clarifying principal identity and accountable ownership," open under basilpuglisi.	1790192074234_image.png	88ceca007f5c8a779ab14310affad1a8cb08707e9c2ac664092b1fb12b19b11f
70	Issue #70 (Comment 6), "Source custody for research-driven requirements and crosswalk citations," open under basilpuglisi.	1790192156200_image.png	71e92352ce51d9a1a6a0c0747927e694992389264e62e1ce781c421f53cf5088
71	Issue #71 (Comment 3), "Approval fatigue: possible pattern-level signals for partial coverage," open under basilpuglisi.	1790192198684_image.png	10ff84f45ec7cd5425495d03d85a926cc8616551f18ee6d1b25694bb3366ccba
72	Issue #72 (Comment 7), "Issue 13: anchoring intervals, history-rewrite risk, and halt authority," open under basilpuglisi.	1790192317216_image.png	a4995cf85c3c0a7cdbaa6cb593ea2b00458d8dc45c36cdd0c48ca9450d829abf
73	Issue #73 (Comment 5), "Evidence-basis declarations and classification of human decisions," open under basilpuglisi.	1790192413337_image.png	90844e5611d199617cd2f525cb056b6edb6b2315b9a8362398b0cae389f9bb3
74	Issue #74 (Comment 8), "H.R. 10362 alignment: per-transaction inference evidence and a NIST-facing crosswalk," open under basilpuglisi.	1790192475417_image.png	1ab4fe2e918e405058974a45c888c62b42b30ad7c30be1231f89e6ff1d511cb2
75	Issue #75 (Comment 9), "Recording who accepted the control set that evidence binds to," open under basilpuglisi.	1790192518177_image.png	de635808108e0968aa6514951c18913e86602ae45fe2f5d855f095be2e9fa164
B	Issue list after seven filings (#67 through #72 visible), confirming sequence before Comment 5 was filed.	1790192357097_image.png	3e16f7fdb34bcd86fb175b22013f034f0c0a1ead8712ceda64b70787d713c287
C	Issue list after all nine filings, showing #67 through #75 opened by basilpuglisi.	1790192550969_image.png	6f5db9e2e983dec6154242d0611c00ca01bb9ea25f4e49eb7249f874558e42a7

Exhibit A

Prefilled issue form for Comment 1 before submission, showing the title and body as drafted. Source file 1790191857147_image.png.

Proof

paste

Proof

Proof

N

Public

AI Go

Proof

Kimi

Mini

Meta

Open

Qwen

Micro

Work

What

Quote

Who

HQ Pr

Ask Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/new?title=C8+open+item+2%3A+clarifying+human+authority+at+Tier+4+and+the+root+of+trust&body=%2A%2ADraft+reference%3A%2A%2A+C8%2C+open+item+2+%28Tier+4+an...

Bookmarks

(3) Create - Roblox

My CJ

MARBLE RUN

Adobe Acrobat

All Bookmarks

Create new issue

Add a title *

C8 open item 2: clarifying human authority at Tier 4 and the root of trust

Add a description

Write

Preview

H B I

Draft reference: C8, open item 2 ("Tier 4 and human authority"); C8 tier table, Accountability row

The Tier 4 accountability cell names "an accountable person whose authorization is un-bypassable," while the tier ladder grades who must be trusted to believe the evidence. These are distinct questions, yet the current table places the accountability question inside a scale designed mainly to measure evidentiary trust. The accountability row itself names a human at Tier 2 ("Named human and institution, bound to policy"), not only at Tier 4. That suggests accountability runs across the tiers rather than rising with them. As the working group resolves this item, it may help to settle where human authority belongs: the Tier 4 root-of-trust cell, or a separate axis the evidence references.

Related Checkpoint-Based Governance work published in September 2025 places a designated human at defined decision points. It treats any inability to show who decided, on what criteria, and with what rationale as an evidential gap. Under that framework, the human at each checkpoint decides to accept, modify, or reject the output. The record shows who decided, on what evidence, why, and what dissent ran against it.

Related agent-architecture work from February 2026 separates agent automation, which it names Responsible AI, from AI Governance, in which a named human holds binding decision authority over the output. A June 2026 paper develops that distinction into two axes, process control and accountability, and holds that only the accountability axis determines governance.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). "Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration" (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>

- Updated: "Checkpoint-Based Governance: Presence is not authority" (current version). <https://basilpuglisi.com/cbg/>

- Puglisi, B. C. (2026, February 3). "HAIA-RECLIN governance architecture specification" (v1.7), Section 2, Three HAIA operating models. GitHub. <https://github.com/basilpuglisi/HAIA/commit/d8cc73a05fb47b491b96fdb7b3fb435188fd80f>

- Published article: "HAIA-RECLIN: Agent governance architecture, an audit-grade multi-AI collaboration, EU regulatory compliance edition" (February 6, 2026). <https://basilpuglisi.com/haia-reclin-agent-governance-architecture-an-audit-grade-multi-ai-collaboration-for-eu-regulatory-compliance-edition/>

- Puglisi, B. C. (2026, June 4). "Why agentic AI was always going to fail". <https://basilpuglisi.com/why-agentic-ai-was-always-going-to-fail/>

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

Assignees

No one assigned

Labels

No labels

Projects

No projects

Milestone

No milestone

☐ Please follow this repository's [contributing guidelines](#) and [security policy](#).

☐ Create more

Cancel

Create

Exhibit 67

Issue #67 (Comment 1), "C8 open item 2: clarifying human authority at Tier 4 and the root of trust," open under basilpuglisi. Source file 1790191921653_image.png.

LFDT-ProofOfControl / ov-poc-standard

Type 17 to search

CodeIssues 13Pull requests 1AgentsActionsProjectsSecurity and qualityInsights

C8 open item 2: clarifying human authority at Tier 4 and the root of trust #67

Open

basilpuglisi opened 1m ago

Draft reference: C8, open item 2 ("Tier 4 and human authority"); C8 tier table, Accountability row

The Tier 4 accountability cell names "an accountable person whose authorization is un-bypassable," while the tier ladder grades who must be trusted to believe the evidence. These are distinct questions, yet the current table places the accountability question inside a scale designed mainly to measure evidentiary trust. The accountability row itself names a human at Tier 2 ("Named human and institution, bound to policy"), not only at Tier 4. That suggests accountability runs across the tiers rather than rising with them. As the working group resolves this item, it may help to settle where human authority belongs: the Tier 4 root-of-trust cell, or a separate axis the evidence references.

Related Checkpoint-Based Governance work published in September 2025 places a designated human at defined decision points. It treats any inability to show who decided, on what criteria, and with what rationale as an evidential gap. Under that framework, the human at each checkpoint decides to accept, modify, or reject the output. The record shows who decided, on what evidence, why, and what dissent ran against it.

Related agent-architecture work from February 2026 separates agent automation, which it names Responsible AI, from AI Governance, in which a named human holds binding decision authority over the output. A June 2026 paper develops that distinction into two axes, process control and accountability, and holds that only the accountability axis determines governance.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). *Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration* (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>
 - Updated: *Checkpoint-Based Governance: Presence is not authority* (current version). <https://basilpuglisi.com/cbg/>
- Puglisi, B. C. (2026, February 3). *HAIA-RECCLIN governance architecture specification* (v1.7), Section 2, Three HAIA operating models. GitHub. [basilpuglisi/HAIA@d8cc73a](https://github.com/basilpuglisi/HAIA@d8cc73a)
 - Published article: *HAIA-RECCLIN: Agent governance architecture, an audit-grade multi-AI collaboration, EU regulatory compliance edition* (February 6, 2026). <https://basilpuglisi.com/haia-recclin-agent-governance-architecture-an-audit-grade-multi-ai-collaboration-for-eu-regulatory-compliance-edition/>
- Puglisi, B. C. (2026, June 4). *Why agentic AI was always going to fail*. <https://basilpuglisi.com/why-agentic-ai-was-always-going-to-fail/>

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications

Customize

Subscribe

You're not receiving notifications from this thread.

Participants

Issue #68 (Comment 2), "Clarifying the relationship between open verification and human responsibility," open under basilpuglisi. Source file 1790192024034_image.png.

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/68

All Bookmarks

LFDT-ProofOfControl / ov-poc-standard

Type / to search

<> Code

Issues 14

Pull requests 1

Agents

Actions

Projects

Security and quality

Insights

Clarifying the relationship between open verification and human responsibility #68

New issue

basilpuglisi opened 1m ago

Draft reference: Google Doc Section 2, "Why open verification must replace human verification"; docs/why-verification-matters.md; C8 Accountability row; C7.5

The informative case says that open verification "must replace human verification" (Google Doc, Section 2 heading). It adds that "the answer cannot be more human review" (docs/why-verification-matters.md). C8 likewise describes Tiers 1 and 2 as human verification and Tiers 3 and 4 as mechanical. The normative text, however, keeps "an accountable person whose authorization is un-bypassable" at Tier 4. C7.5 also states that validation "stays a human responsibility." A reader could therefore take the informative framing more broadly than intended, even though the normative requirements preserve human responsibility for validation and accountability.

Aligning the informative case with the normative core could reduce that ambiguity. One option separates machine verification of what an agent executed from human authority over the controls it executed against, which matches the draft's own line between verification and validation. Related policy work from March 2026 treats verification infrastructure as serving governance, so that human checkpoint decisions rest on verifiable evidence rather than contractual assurance. The request concerns the framing only and does not ask for human review of machine-speed execution evidence.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). *Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration* (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>
 - Updated: *Checkpoint-Based Governance: Presence is not authority* (current version). <https://basilpuglisi.com/cbg/>
- Puglisi, B. C. (2026, February 3). *HAIA-RECCLIN governance architecture specification* (v1.7), Section 2, Three HAIA operating models. GitHub. [basilpuglisi/HAIA@ d8cc73a](https://github.com/basilpuglisi/HAIA@d8cc73a)
 - Published article: *HAIA-RECCLIN: Agent governance architecture, an audit-grade multi-AI collaboration, EU regulatory compliance edition* (February 6, 2026). <https://basilpuglisi.com/haia-recclin-agent-governance-architecture-an-audit-grade-multi-ai-collaboration-for-eu-regulatory-compliance-edition/>
- Puglisi, B. C. (2026, March 6). *Verified AI Inference Standards Act: A congressional framework for trustworthy AI data processing* (AI Provider Plurality Congressional Package, Document 5 of 5), Appendix B. <https://basilpuglisi.com/wp-content/uploads/2026/03/Verified-AI-Inference-Standards-Act-or-VAISA.pdf>

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications

Customize

Unsubscribe

You're receiving notifications because you're subscribed to this thread.

Participants

Issue #69 (Comment 4), "Clarifying principal identity and accountable ownership," open under basilpuglisi. Source file 1790192074234_image.png.

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/69

★ Bookmarks

(3) Create - Roblox

My CJ

MARBLE RUN

Adobe Acrobat

All Bookmarks

LFDTProofOfControl / ov-poc-standard

Type / to search

+ - ↺ 🔗 📄 📧 👤

<> Code Issues 15 Pull requests 1 Agents Actions Projects Security and quality Insights

Clarifying principal identity and accountable ownership #69

New issue ✎

🔖 Open

basilpuglisi opened 1m ago

⋮

Draft reference: C5.1.1; C4.2.2

C5.1.1 resolves every sampled action to a named principal, and C4.2.2 validates each delegation back to the originating principal. Neither appears intended to establish who ultimately answers for the outcome. An authenticated principal may still not be the person who owns the result or supervises the agent. The same principal may also not hold the relevant organizational or professional responsibility, or be the party that may bear legal responsibility for the outcome.

The working group may want a distinct accountable-person field, using the term the Tier 4 accountability cell already uses. A plain statement that principal identity does not, by itself, assign accountability would serve the same purpose. Without one of the two, audits and disputes could treat a verified principal as the responsible party.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>
 - Updated: Checkpoint-Based Governance: Presence is not authority (current version). <https://basilpuglisi.com/cbg/>
- Puglisi, B. C. (2026, June 4). Why agentic AI was always going to fail. <https://basilpuglisi.com/why-agentic-ai-was-always-going-to-fail/>

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

🔒

Add a comment

WritePreview

HBI≡<>🔗≡≡≡@📧📄↶

Use Markdown to format your comment

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

NotificationsCustomize

Unsubscribe

You're receiving notifications because you're subscribed to this thread.

Participants

Issue #70 (Comment 6), "Source custody for research-driven requirements and crosswalk citations," open under basilpuglisi. Source file 1790192156200_image.png.

Proof

paste

Proof

Sr

Public

AI Go

Proof

Kimi

MiniM

Meta

Open

Qwen

Micro

Work

What

Quote

Who's

HQ P

+

Ask Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/70

🔖 ☆ 📁 | All Bookmarks

☰ LFDT-ProofOfControl / ov-poc-standard ▾ 🔍 Type ↵ to search 🗑️ + ⌚ ⚙️ 📄 📧 👤

<> Code Issues 16 Pull requests 1 Agents Actions Projects Security and quality Insights

Source custody for research-driven requirements and crosswalk citations #70 ✎

New issue 📋

 basilpuglisi opened 1m ago ...

Draft reference: Appendix D, issue 12; README mapping caveat; mappings/

Appendix D issue 12 leaves open whether the 2026 research citations verify against the primary literature. Separately, the independent mapping review found 96 clause-citation errors. It also found that some framework mappings rely heavily on label-level or chapter-level citations, and that all 60 CSA AARM citations were label-level rather than clause-level. Before ratification, reviewers could confirm the basis of each research-driven requirement and crosswalk row if a record showed what was verified for each cited source. That record would note when the source was checked, what it said, and whether it supported the requirement.

Published source-custody work offers one documented form for that record. It is an author-side record, so the draft would still need a public, comparable disclosure for research-driven requirements.

Published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2026, June 2). *SCOPE: Source custody observable publication evidence*. <https://basilpuglisi.com/scope-source-custody-observable-publication-evidence/>
 - Updated: current version at the same address.
- Puglisi, B. C. (2026, June 3). *Fault-based publication ethics: The case for source custody in an era of AI citation contamination*. <https://basilpuglisi.com/fault-based-publication-ethics-source-custody/> (SSRN Abstract 6872038).

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

😊

Add a comment

Write Preview

H B I ☰ <> 🔗 ⌨️ ⌨️ ⌨️ @ 📎 📄 ↶

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications Customize

Subscribe 🔔

You're not receiving notifications from this thread.

Participants

Exhibit 71

Issue #71 (Comment 3), "Approval fatigue: possible pattern-level signals for partial coverage," open under basilpuglisi. Source file 1790192198684_image.png.

Proof

paste

Proof

Proof

AI

Public

AI Go

Proof

Kimi

Mini

Meta

Open

Qwen

Micro

Works

What

Quote

Who

HQ Pr

Ask Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/71

Bookmarks (3) Create - Roblox My CJ MARBLE RUN Adobe Acrobat

LFDT-ProofOfControl / ov-poc-standard

Type 71 to search

+ -

<> Code Issues 17 Pull requests 1 Agents Actions Projects Security and quality Insights

Approval fatigue: possible pattern-level signals for partial coverage #71

Open

basilpuglisi

opened 1m ago

Draft reference: Appendix C, threat "Human-agent trust exploitation / approval fatigue" (coverage Partial); docs/use-cases/COVERAGE.md (Human oversight and disclosure, 0 of 2); C4.1.6

Appendix C rates approval fatigue as Partial because the evidence captures what was presented and decided, not the fatigue itself. The use-case coverage table also shows no worked example for the human oversight threat group. On their own, the approval records required by C4.1.6 do not appear sufficient to define a reliable rubber-stamp signal. C4.1.6 specifies no denominator of approval opportunities, and it does not distinguish rejection, deferral, and later reversal of the same decision.

Appendix B already notes that "human oversight quality is outside the scope of cryptographic proof." A pattern-level measure could stay focused on observable review behavior across records rather than trying to infer a reviewer's state of mind. Related work published in September 2025 names automation bias drift, where reviewers progressively defer to AI recommendations without critical evaluation, as a core governance failure. Later work proposes detection metrics for it.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). *Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration* (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>
 - Updated: *Checkpoint-Based Governance: Presence is not authority* (current version). <https://basilpuglisi.com/cbg/>
- Puglisi, B. C. (2026, February 4). *HAIA-RECCLIN agent architecture specification: EU compliance version* (Version 2.2), Appendix B, Automation bias detection metric definitions. GitHub. [@ 6378a4c](https://github.com/basilpuglisi/HAIA)

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications

Customize

Unsubscribe

You're receiving notifications because you're subscribed to this thread.

Participants

Add a comment

Issue #73 (Comment 5), "Evidence-basis declarations and classification of human decisions," open under basilpuglisi. Source file 1790192413337_image.png.

ProofpasteProofPublicAI GoKimiMiniMetaOpenQwenMicroWorkWhatQuoteWho'sHQ PAsk Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/73

Bookmarks(3) Create - RobloxMy CIMARBLE RUNAdobe AcrobatAll Bookmarks

LFD-ProofOfControl / ov-poc-standard

Type 77 to search

+--@

<> CodeIssues 19Pull requests 1AgentsActionsProjectsSecurity and qualityInsights

Evidence-basis declarations and classification of human decisions #73

New issue

basilpuglisi opened 1m ago

Draft reference: C4.1.6; C10.1.4 and C10.2; the declared-control-set requirement described on the Proof-of-Control page as in drafting

There is published prior work relevant to two open threads. The declared-control-set requirement and the C10 conformance statement both ask an operator to declare what a claim rests on. CARCS, published April 23, 2026, proposes that each governance record declare its evidence status class before generation. It also proposes an evidence ceiling rule, so the record claims no more than its evidence supports. CARCS distinguishes Attestation Grade from Hash Verified records and treats a locally computed hash as self-attested integrity. External verification, in its terms, comes from anchoring the hash to a third-party timestamp service or equivalent infrastructure.

C4.1.6 records that a human approved, not what kind of engagement occurred. CARCS proposes classifying the human act as corrective override, creative supersession, checkpoint confirmation, or deferred decision. It pairs that classification with the decision rationale and proposes that dissent between sources be preserved rather than averaged away. These proposals remain under the cited work's license and are referenced here as proposals, not offered as replacement text.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2025, September 23). *Checkpoint-Based Governance: An implementation framework for accountable human-AI collaboration* (v2 drafting). <https://basilpuglisi.com/checkpoint-based-governance-an-implementation-framework-for-accountable-human-ai-collaboration-v2-drafting/>
 - Updated: *Checkpoint-Based Governance: Presence is not authority* (current version). <https://basilpuglisi.com/cbg/>
- Puglisi, B. C. (2026, April 23). *HAIA-CARCS: Compliance accountability record and case study* (v1.33). GitHub. [basilpuglisi/HAIA@7d5076a](#)
 - Updated: *CARCS: Compliance accountability record and case study* (current version). <https://basilpuglisi.com/haia-carcs-compliance-accountability-record-case-study/>

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

AssigneesNo one assigned

LabelsNo labels

TypeNo type

ProjectsNo projects

MilestoneNo milestone

RelationshipsNone yet

DevelopmentNo branches or pull requests

NotificationsCustomizeSubscribeYou're not receiving notifications from this thread.

Participants

Exhibit 74

Issue #74 (Comment 8), "H.R. 10362 alignment: per-transaction inference evidence and a NIST-facing crosswalk," open under basilpuglisi. Source file 1790192475417_image.png.

Proof

paste

Proof

Proof

H

Public

AI Go

Proof

Kimi

MiniM

Meta

Open

Qwen

Micro

Works

What

Quote

Who's

HQ Pr

+

Ask Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues/74

Bookmarks (3) Create - Roblox My CJ MARBLE RUN Adobe Acrobat

All Bookmarks

LFDT-ProofOfControl / ov-poc-standard

Type 7 to search

+ -

🔍

📄

📧

👤

<> Code

🕒 Issues 20

🔗 Pull requests 1

👤 Agents

🔄 Actions

📁 Projects

🛡️ Security and quality

🔍 Insights

H.R. 10362 alignment: per-transaction inference evidence and a NIST-facing crosswalk #74

New issue

🔗 Open

basilpuglisi

opened 1m ago

Draft reference: Launch materials citing H.R. 10362 (Stop Rogue AI Act); C1 Provenance; C7.7 Interoperable property; C8.1.7

The launch materials present Proof-of-Control as written for the question H.R. 10362 asks: verifiable agent identity and provenance, not a vendor's word. The introduced bill directs NIST to set those standards. It also asks for an agent inventory, identity verification at network and application layers, and the power to allow, deny, or revoke agent access. The draft covers the evidence part of that list. A short informative crosswalk could map the bill's provisions to the draft's requirements, marking which ones the draft evidences and which it leaves to enforcement or discovery. That would help NIST and other readers see the fit without overstating it.

Related policy work published in March 2026 proposes a comparable NIST directive at the inference boundary. It proposes a signed processing receipt that binds each inference transaction to an attested environment through a nonce and an input hash. It also proposes that NIST define per-request evidence schemas and harmonize attestation formats across hardware vendors. That work relies on vendor-rooted hardware attestation, which the draft places at Tier 2 unless composed with independent anchoring (C8.1.7). It is offered as related prior work on the evidence artifact, not as a Tier 3 mechanism. Where attestation is unavailable, the same work proposes a gate at which a named human reviews the dispatch and decides whether it proceeds, with that decision logged.

Prior published work, offered as reference under its existing license and not as contributed text:

- Puglisi, B. C. (2026, March 6). *Verified AI Inference Standards Act: A congressional framework for trustworthy AI data processing* (AI Provider Plurality Congressional Package, Document 5 of 5). <https://basilpuglisi.com/wp-content/uploads/2026/03/Verified-AI-Inference-Standards-Act-or-VAISA.pdf>

Basil C. Puglisi, MPA. Independent practitioner, basilpuglisi.com.

👤

Assignees

No one assigned

Labels

No labels

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications

Customize

🔔 Unsubscribe

You're receiving notifications because you're subscribed to this thread.

Participants

Add a comment

Exhibit B

Issue list after seven filings (#67 through #72 visible), confirming sequence before Comment 5 was filed. Source file 1790192357097_image.png.

LFDT-ProofOfControl / ov-poc-standard

Type 7 to search

+ - ↺ 🔍 📄 📧 👤

<> Code

Issues 18

Pull requests 1

Agents

Actions

Projects

Security and quality

Insights

Issues

Assigned to me

Created by me

Mentioned

Recent activity

Views

Projects

Milestones

Labels

is:issue state:open

Open 19 Closed 2

Author Labels Projects Milestones Assignees Types ↕ Best match

Issue 13: anchoring intervals, history-rewrite risk, and halt authority

LFDT-ProofOfControl/ov-poc-standard#72 · basilpuglisi opened 1m ago

Approval fatigue: possible pattern-level signals for partial coverage

LFDT-ProofOfControl/ov-poc-standard#71 · basilpuglisi opened 2m ago

Source custody for research-driven requirements and crosswalk citations

LFDT-ProofOfControl/ov-poc-standard#70 · basilpuglisi opened 3m ago

Clarifying principal identity and accountable ownership

LFDT-ProofOfControl/ov-poc-standard#69 · basilpuglisi opened 4m ago

Clarifying the relationship between open verification and human responsibility

LFDT-ProofOfControl/ov-poc-standard#68 · basilpuglisi opened 5m ago

C8 open item 2: clarifying human authority at Tier 4 and the root of trust

LFDT-ProofOfControl/ov-poc-standard#67 · basilpuglisi opened 7m ago

How should a token represent an evaluation that could not complete?

LFDT-ProofOfControl/ov-poc-standard#66 · rakeshghel01 opened 1d ago

Is a declared source that stops producing evidence meant to be detectable below Level 4?

LFDT-ProofOfControl/ov-poc-standard#65 · rakeshghel01 opened 1d ago

Field note: signing-key custody, the Tier 2/3 line, and economic recourse from a live receipt-attestation system

LFDT-ProofOfControl/ov-poc-standard#64 · aglamadrid19 opened 3d ago

Implementation feedback: independent appraisal of a Proof-of-Control reference token

LFDT-ProofOfControl/ov-poc-standard#63 · terryncw opened 5d ago

Exhibit C

Issue list after all nine filings, showing #67 through #75 opened by basilpuglisi. Source file 1790192550969_image.png.

ProofpasteProofProofIs xPublicAI GoProofKimiMiniMetaOpenQwenMicroWorksWhatsAppQuoteWho'sHQ Pr+Ask Gemini

github.com/LFDT-ProofOfControl/ov-poc-standard/issues

Bookmarks(3) Create - RobloxMy CMMARBLE RUNAdobe Acrobat

LFDT-ProofOfControl / ov-poc-standard

Type to search

+ -

🔍

📄

📧

👤

<> Code

🕒 Issues 21

🔗 Pull requests 1

🛠 Agents

🔄 Actions

📁 Projects

🔒 Security and quality

📊 Insights

Issues

Assigned to me

Created by me

Mentioned

Recent activity

Views

Projects

Milestones

Labels

Feedback

Preview

Collapse sidebar

Issues

New issue

is:issue state:open

Open 22Closed 2

AuthorLabelsProjectsMilestonesAssigneesTypesBest match

Recording who accepted the control set that evidence binds to

LFDT-ProofOfControl/ov-poc-standard#75 · basilpuglisi opened 1m ago

H.R. 10362 alignment: per-transaction inference evidence and a NIST-facing crosswalk

LFDT-ProofOfControl/ov-poc-standard#74 · basilpuglisi opened 1m ago

Evidence-basis declarations and classification of human decisions

LFDT-ProofOfControl/ov-poc-standard#73 · basilpuglisi opened 2m ago

Issue 13: anchoring intervals, history-rewrite risk, and halt authority

LFDT-ProofOfControl/ov-poc-standard#72 · basilpuglisi opened 3m ago

Approval fatigue: possible pattern-level signals for partial coverage

LFDT-ProofOfControl/ov-poc-standard#71 · basilpuglisi opened 5m ago

Source custody for research-driven requirements and crosswalk citations

LFDT-ProofOfControl/ov-poc-standard#70 · basilpuglisi opened 6m ago

Clarifying principal identity and accountable ownership

LFDT-ProofOfControl/ov-poc-standard#69 · basilpuglisi opened 7m ago

Clarifying the relationship between open verification and human responsibility

LFDT-ProofOfControl/ov-poc-standard#68 · basilpuglisi opened 8m ago

C8 open item 2: clarifying human authority at Tier 4 and the root of trust

LFDT-ProofOfControl/ov-poc-standard#67 · basilpuglisi opened 10m ago

How should a token represent an evaluation that could not complete?

LFDT-ProofOfControl/ov-poc-standard#66 · rakeshgohel01 opened 1d ago

Record notes

The issues are public and permanent on GitHub under the filer's account. The screenshots record each issue as it appeared at filing; the live issues remain the primary record. Cited works are identified by first public introduction, with the current version listed beneath where one exists, and GitHub commit hashes carry the dates where a web page has been modified since first publication.

#Assisted using the HAIA Ecosystem | CC BY-NC-SA 4.0 Free for personal, educational, and noncommercial research use with attribution. Commercial exploitation, paid productization, and enterprise commercialization require separate permission and licensing.